

Andrii Lunin

Professor Mohamed Zahran

Parallel Processing

Lab 1

November 11, 2020

Lab 1

- 1) Once the module is loaded with `module load mpi/openmpi-x86_64`, you can run my program.
- 2) I have provided a handy [Makefile](#), for compiling.
 - 1) On `$make` it compiles the `genprimes.c` file into the executive `genprimes`.
 - 2) On `$make clean`, it **removes** the executive file **AND** all of the `.txt` files. Be mindful of that if you run it in some other folder apart from the folder with my code.
- 3) For executing the program run `$mpiexec -n <P> genprimes <N>`, where P stands for the number of processes, and N stands for the upper bound on the range of primes.
- 4) I have to note, that I tailored my program to be as *quick* and as *efficient* as possible for large numbers. I have explained my algorithm in detail in comments in the C file, but I wanted to make sure to assert that the best efficiency comes with $P > 10M$. Once you consider $> 1B$, the speedup for a 100 processes is massive (I have provided the times for $p = 1B$ for 10 and 100 processes in the table below).

Time

#P vs N	1,000	10,000	100,000	1,000,000	10,000,000	Extra: 1B
1	0.249s	0.250s	0.273s	0.701s	10.946s	X
2	0.265s	0.265s	0.273s	0.457s	4.485s	X
5	0.298s	0.296s	0.301s	0.380s	1.645s	X
10	0.370s	0.372s	0.377s	0.417s	1.023s	289.819s
100	3.721s	3.766s	3.703s	3.678s	4.072s	42.161s

Speedup

#P vs N	1,000	10,000	100,000	1,000,000	10,000,000
1	1	1	1	1	1
2	0.939	0.943	1	1.534	2.44
5	0.836	0.845	0.907	1.845	6.654
10	0.839	0.672	0.724	1.681	10.7
100	0.067	0.066	0.074	0.190	2.688

Efficiency

#P vs N	1,000	10,000	100,000	1,000,000	10,000,000
1	1	1	1	1	1
2	0.4695	0.4715	0.5	0.767	1.22
5	0.1672	0.169	0.1814	0.369	1.3308
10	0.0839	0.0672	0.0724	0.1681	1.07
100	0.00067	0.00066	0.00074	0.0019	0.02688

*Note: all times provided are averaged over 10 consecutive executions.